

# Project\_4\_Lung\_cancer\_detection

February 27, 2025

```
[1]: #load required library
import numpy as np, pandas as pd, matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: #load dataest
df = pd.read_csv('dataset_med.csv')
df.head()
```

```
[2]:
```

|   | id | age | gender | country     | diagnosis_date | cancer_stage | family_history | \ |
|---|----|-----|--------|-------------|----------------|--------------|----------------|---|
| 0 | 1  | 64  | Male   | Sweden      | 05 04 2016     | Stage I      | Yes            |   |
| 1 | 2  | 50  | Female | Netherlands | 20 04 2023     | Stage III    | Yes            |   |
| 2 | 3  | 65  | Female | Hungary     | 05 04 2023     | Stage III    | Yes            |   |
| 3 | 4  | 51  | Female | Belgium     | 05 02 2016     | Stage I      | No             |   |
| 4 | 5  | 37  | Male   | Luxembourg  | 29 11 2023     | Stage I      | No             |   |

|   | smoking_status | bmi  | cholesterol_level | hypertension | asthma | cirrhosis | \ |
|---|----------------|------|-------------------|--------------|--------|-----------|---|
| 0 | Passive Smoker | 29.4 | 199               | 0            | 0      | 1         |   |
| 1 | Passive Smoker | 41.2 | 280               | 1            | 1      | 0         |   |
| 2 | Former Smoker  | 44.0 | 268               | 1            | 1      | 0         |   |
| 3 | Passive Smoker | 43.0 | 241               | 1            | 1      | 0         |   |
| 4 | Passive Smoker | 19.7 | 178               | 0            | 0      | 0         |   |

|   | other_cancer | treatment_type | end_treatment_date | survived |
|---|--------------|----------------|--------------------|----------|
| 0 | 0            | Chemotherapy   | 10 09 2017         | 0        |
| 1 | 0            | Surgery        | 17 06 2024         | 1        |
| 2 | 0            | Combined       | 09 04 2024         | 0        |
| 3 | 0            | Chemotherapy   | 23 04 2017         | 0        |
| 4 | 0            | Combined       | 08 01 2025         | 0        |

```
[3]: #remove unused columns
df.drop(columns=['id', 'end_treatment_date', 'diagnosis_date'], inplace=True)
```

```
[4]: #check shape
df.shape
```

[4]: (890000, 14)

```
[5]: #check statistical report
df.describe()
```

```
[5]:
```

|       | age           | bmi           | cholesterol_level | hypertension  | \ |
|-------|---------------|---------------|-------------------|---------------|---|
| count | 890000.000000 | 890000.000000 | 890000.000000     | 890000.000000 |   |
| mean  | 55.007008     | 30.494172     | 233.633916        | 0.750024      |   |
| std   | 9.994485      | 8.368539      | 43.432278         | 0.432999      |   |
| min   | 4.000000      | 16.000000     | 150.000000        | 0.000000      |   |
| 25%   | 48.000000     | 23.300000     | 196.000000        | 1.000000      |   |
| 50%   | 55.000000     | 30.500000     | 242.000000        | 1.000000      |   |
| 75%   | 62.000000     | 37.700000     | 271.000000        | 1.000000      |   |
| max   | 104.000000    | 45.000000     | 300.000000        | 1.000000      |   |

|       | asthma        | cirrhosis     | other_cancer  | survived      |
|-------|---------------|---------------|---------------|---------------|
| count | 890000.000000 | 890000.000000 | 890000.000000 | 890000.000000 |
| mean  | 0.469740      | 0.225956      | 0.088157      | 0.220229      |
| std   | 0.499084      | 0.418211      | 0.283524      | 0.414401      |
| min   | 0.000000      | 0.000000      | 0.000000      | 0.000000      |
| 25%   | 0.000000      | 0.000000      | 0.000000      | 0.000000      |
| 50%   | 0.000000      | 0.000000      | 0.000000      | 0.000000      |
| 75%   | 1.000000      | 0.000000      | 0.000000      | 0.000000      |
| max   | 1.000000      | 1.000000      | 1.000000      | 1.000000      |

```
[6]: #check information about each column
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 890000 entries, 0 to 889999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   890000 non-null  int64
1   gender                890000 non-null  object
2   country               890000 non-null  object
3   cancer_stage          890000 non-null  object
4   family_history        890000 non-null  object
5   smoking_status        890000 non-null  object
6   bmi                   890000 non-null  float64
7   cholesterol_level     890000 non-null  int64
8   hypertension          890000 non-null  int64
9   asthma                890000 non-null  int64
10  cirrhosis             890000 non-null  int64
11  other_cancer           890000 non-null  int64
12  treatment_type         890000 non-null  object
13  survived               890000 non-null  int64
dtypes: float64(1), int64(7), object(6)
```

memory usage: 95.1+ MB

```
[7]: #checking numerical column
df.select_dtypes(include=[int,float]).columns
```

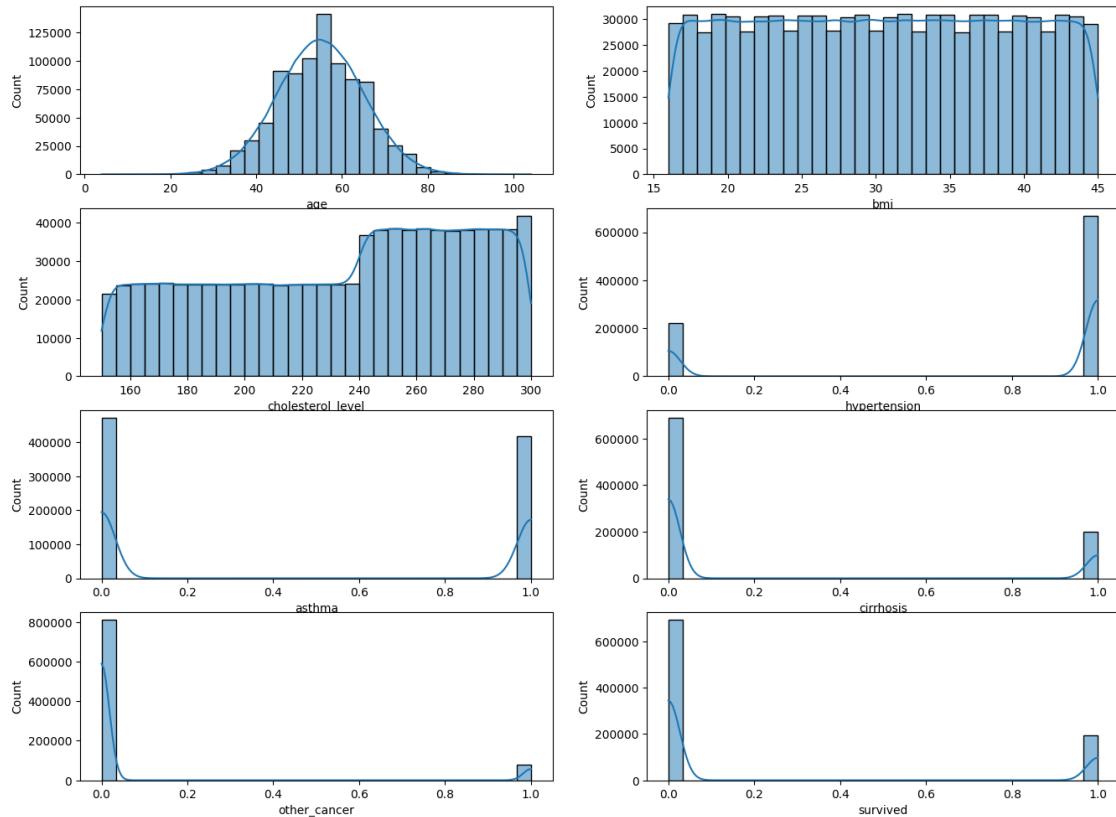
```
[7]: Index(['age', 'bmi', 'cholesterol_level', 'hypertension', 'asthma',
         'cirrhosis', 'other_cancer', 'survived'],
         dtype='object')
```

```
[8]: #checking categorical columns
df.select_dtypes(include=object).columns
```

```
[8]: Index(['gender', 'country', 'cancer_stage', 'family_history', 'smoking_status',
         'treatment_type'],
         dtype='object')
```

```
[9]: #find numerical column
num_col = ['age', 'bmi', 'cholesterol_level', 'hypertension', 'asthma',
          ↪ 'cirrhosis', 'other_cancer', 'survived']
```

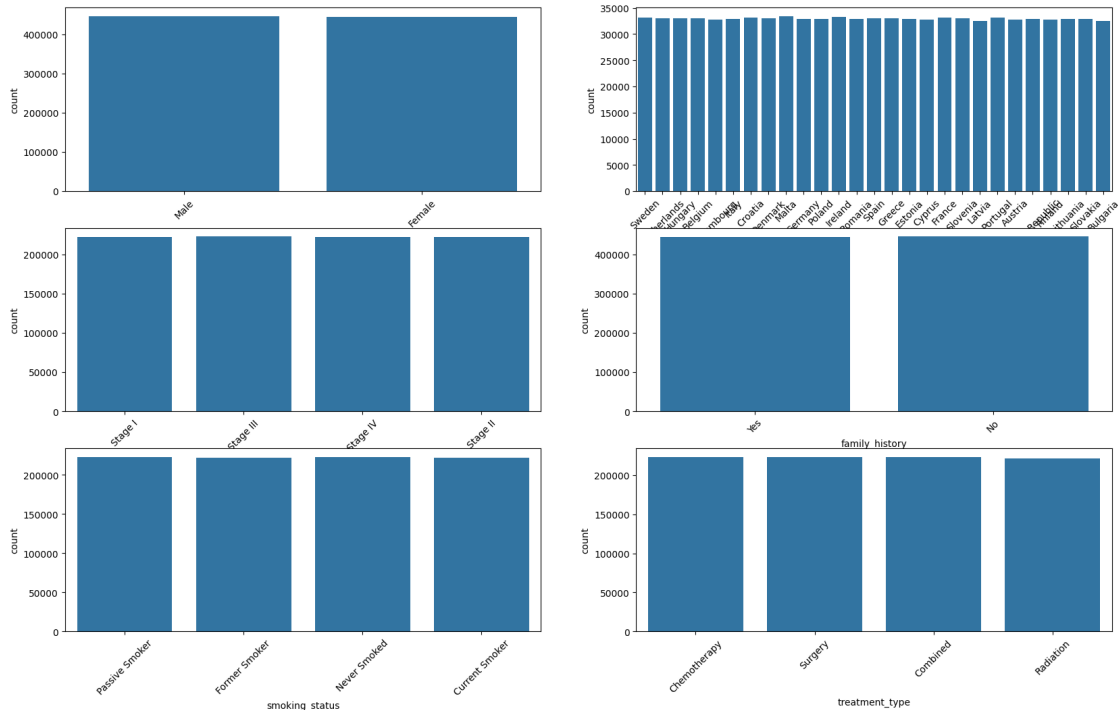
```
[10]: #Data distribution
plt.figure(figsize=(16, 12))
for i, col in enumerate(num_col):
    plt.subplot(4,2,i+1)
    sns.histplot(df[col], bins=30, kde=True)
```



above all Numerical columns distribution is correct and range in thier scaled.

But Survived columns having Imbalanced.

```
[11]: #categorical columns distribution
col = ['gender', 'country', 'cancer_stage', 'family_history', 'smoking_status', 'treatment_type']
plt.figure(figsize=(20, 12))
for i, col in enumerate(col):
    plt.subplot(3,2,i+1)
    sns.countplot(x=df[col])
    plt.xticks(rotation=45)
```



```
[12]: #find categorical columns
df.select_dtypes(include=object).columns
```

```
[12]: Index(['gender', 'country', 'cancer_stage', 'family_history', 'smoking_status',
          'treatment_type'],
          dtype='object')
```

```
[13]: #make list of categorical columns
cat_col = ['gender', 'country', 'cancer_stage', 'family_history',
          ↪ 'smoking_status', 'treatment_type']
```

```
[14]: #load onehot_encoder and fit in categorical col
onehot_encoder = OneHotEncoder(drop='first', sparse_output=False)
onehot_encoder.fit(df[cat_col])

#transform the onehotencoder to the dataset
encoded_col = onehot_encoder.transform(df[cat_col])
encoded_df = pd.DataFrame(encoded_col, columns=onehot_encoder.
    ↪ get_feature_names_out(cat_col))
df_encoded = pd.concat([df.drop(columns=cat_col), encoded_df], axis=1)
df_encoded.shape
```

```
[14]: (890000, 14)
```

```
[15]: #check encoded df
df_encoded.head()
```

```
[15]:  age    bmi  cholesterol_level  hypertension  asthma  cirrhosis  \
0    64   29.4                199                0        0        1
1    50   41.2                280                1        1        0
2    65   44.0                268                1        1        0
3    51   43.0                241                1        1        0
4    37   19.7                178                0        0        0

    other_cancer  survived  gender_Male  country_Belgium  ...  \
0              0         0          1.0             0.0  ...
1              0         1          0.0             0.0  ...
2              0         0          0.0             0.0  ...
3              0         0          0.0             1.0  ...
4              0         0          1.0             0.0  ...

    cancer_stage_Stage II  cancer_stage_Stage III  cancer_stage_Stage IV  \
0                   0.0                   0.0                   0.0
1                   0.0                   1.0                   0.0
2                   0.0                   1.0                   0.0
3                   0.0                   0.0                   0.0
4                   0.0                   0.0                   0.0

    family_history_Yes  smoking_status_Former Smoker  \
0                   1.0                   0.0
1                   1.0                   0.0
2                   1.0                   1.0
3                   0.0                   0.0
4                   0.0                   0.0

    smoking_status_Never Smoked  smoking_status_Passive Smoker  \
0                   0.0                   1.0
1                   0.0                   1.0
2                   0.0                   0.0
3                   0.0                   1.0
4                   0.0                   1.0

    treatment_type_Combined  treatment_type_Radiation  treatment_type_Surgery
0                   0.0                   0.0                   0.0
1                   0.0                   0.0                   1.0
2                   1.0                   0.0                   0.0
3                   0.0                   0.0                   0.0
4                   1.0                   0.0                   0.0
```

```
[5 rows x 45 columns]
```

```
[16]: #checking columns
df_encoded.columns
```

```
[16]: Index(['age', 'bmi', 'cholesterol_level', 'hypertension', 'asthma',
        'cirrhosis', 'other_cancer', 'survived', 'gender_Male',
        'country_Belgium', 'country_Bulgaria', 'country_Croatia',
        'country_Cyprus', 'country_Czech Republic', 'country_Denmark',
        'country_Estonia', 'country_Finland', 'country_France',
        'country_Germany', 'country_Greece', 'country_Hungary',
        'country_Ireland', 'country_Italy', 'country_Latvia',
        'country_Lithuania', 'country_Luxembourg', 'country_Malta',
        'country_Netherlands', 'country_Poland', 'country_Portugal',
        'country_Romania', 'country_Slovakia', 'country_Slovenia',
        'country_Spain', 'country_Sweden', 'cancer_stage_Stage II',
        'cancer_stage_Stage III', 'cancer_stage_Stage IV', 'family_history_Yes',
        'smoking_status_Former Smoker', 'smoking_status_Never Smoked',
        'smoking_status_Passive Smoker', 'treatment_type_Combined',
        'treatment_type_Radiation', 'treatment_type_Surgery'],
        dtype='object')
```

```
[17]: #split data in input and target variable
X = df_encoded.drop(columns=['survived'])
y = df_encoded['survived']

X.shape, y.shape
```

```
[17]: ((890000, 44), (890000,))
```

```
[18]: #loading requieired library, randomforestclassifier and metrix.
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix,
    ↪accuracy_score
```

```
[19]: #split the input feature and target variable in to xtrain, ytrain, xtest and
    ↪y_test
xtrain, xtest, ytrain, ytest = train_test_split(X, y, random_state=42,
    ↪test_size=.2)
xtrain.shape, xtest.shape, ytrain.shape, ytest.shape
```

```
[19]: ((712000, 44), (178000, 44), (712000,), (178000,))
```

```
[20]: #model creating using RF

rfc = RandomForestClassifier(random_state=42)
rfc.fit(xtrain, ytrain)
```

```
[20]: RandomForestClassifier(random_state=42)
```

```
[21]: #predict the model
rfc_pred = rfc.predict(xtest)

#checking classification report and confusion matrix
print(classification_report(ytest, rfc_pred))
print(confusion_matrix(ytest, rfc_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.78      | 1.00   | 0.87     | 138639  |
| 1            | 0.23      | 0.00   | 0.01     | 39361   |
| accuracy     |           |        | 0.78     | 178000  |
| macro avg    | 0.50      | 0.50   | 0.44     | 178000  |
| weighted avg | 0.66      | 0.78   | 0.68     | 178000  |

```
[[137981    658]
 [ 39169    192]]
```

above F1-score for 1 is highly biased that means imbalanced dataset. So, we will remove imbalanced

```
[22]: #load smote library for imbalanced data
from imblearn.over_sampling import SMOTE
sm = SMOTE()
```

```
[23]: #appling SMOTE
X_resamble, y_resamble = sm.fit_resample(X, y)
```

```
[24]: #split dataset into train, test.
xtrain, xtest, ytrain, ytest = train_test_split(X_resamble, y_resamble,
    random_state=42, test_size=.2)
xtrain.shape, xtest.shape, ytrain.shape, ytest.shape
```

```
[24]: ((1110393, 44), (277599, 44), (1110393,), (277599,))
```

```
[25]: #RF model using.
rfc_imbalanced = RandomForestClassifier(random_state=42)
#fitting the dataset.
rfc_imbalanced.fit(xtrain, ytrain)
```

```
[25]: RandomForestClassifier(random_state=42)
```

```
[26]: #predict the model
rfc_imb_pred = rfc_imbalanced.predict(xtest)

#checking the classification and confusion matrix
```



```
print(classification_report(ytest, rfc_imb_pred))
print(confusion_matrix(ytest, rfc_imb_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.78      | 0.97   | 0.87     | 138614  |
| 1            | 0.97      | 0.73   | 0.83     | 138985  |
| accuracy     |           |        | 0.85     | 277599  |
| macro avg    | 0.87      | 0.85   | 0.85     | 277599  |
| weighted avg | 0.87      | 0.85   | 0.85     | 277599  |

```
[[135074  3540]
 [ 38027 100958]]
```

```
[27]: print(accuracy_score(ytest, rfc_imb_pred))
```

```
0.8502624288992395
```

our model has 85% accuracy.

```
[28]: #now save our model.
import joblib

#saving model
joblib.dump(rfc_imbalanced, 'random_forest_model.pkl')

#save onehotencoder
joblib.dump(onehot_encoder, 'onehot_encoder.pkl')
```

```
[28]: ['onehot_encoder.pkl']
```

```
[29]: rfc_imbalanced.feature_importances_
```

```
[29]: array([0.10428272, 0.12043656, 0.11633123, 0.01323727, 0.01100567,
          0.01210045, 0.0094233 , 0.06336654, 0.0063427 , 0.00628626,
          0.00646505, 0.00609679, 0.00604109, 0.00626547, 0.00581283,
          0.00584679, 0.00741864, 0.00632842, 0.00645128, 0.00654287,
          0.00626022, 0.00654555, 0.0060309 , 0.00643853, 0.00615368,
          0.00616094, 0.00571796, 0.00586206, 0.00640378, 0.00598467,
          0.00647707, 0.00628602, 0.00600316, 0.00682076, 0.03497755,
          0.03494464, 0.03319997, 0.06320513, 0.03969594, 0.03628118,
          0.03503372, 0.03706084, 0.03588206, 0.03649175])
```

```
[30]: rfc_imbalanced.feature_names_in_
```

```
[30]: array(['age', 'bmi', 'cholesterol_level', 'hypertension', 'asthma',
          'cirrhosis', 'other_cancer', 'gender_Male', 'country_Belgium',
          'country_Bulgaria', 'country_Croatia', 'country_Cyprus',
```

```
'country_Czech Republic', 'country_Denmark', 'country_Estonia',
'country_Finland', 'country_France', 'country_Germany',
'country_Greece', 'country_Hungary', 'country_Ireland',
'country_Italy', 'country_Latvia', 'country_Lithuania',
'country_Luxembourg', 'country_Malta', 'country_Netherlands',
'country_Poland', 'country_Portugal', 'country_Romania',
'country_Slovakia', 'country_Slovenia', 'country_Spain',
'country_Sweden', 'cancer_stage_Stage II',
'cancer_stage_Stage III', 'cancer_stage_Stage IV',
'family_history_Yes', 'smoking_status_Former Smoker',
'smoking_status_Never Smoked', 'smoking_status_Passive Smoker',
'treatment_type_Combined', 'treatment_type_Radiation',
'treatment_type_Surgery'], dtype=object)
```

```
[31]: #checking the columns
df.columns
```

```
[31]: Index(['age', 'gender', 'country', 'cancer_stage', 'family_history',
'smoking_status', 'bmi', 'cholesterol_level', 'hypertension', 'asthma',
'cirrhosis', 'other_cancer', 'treatment_type', 'survived'],
dtype='object')
```

```
[32]: #checking unique variable in each categorical columns
for i in (df.select_dtypes(include=object)):
    print(f'df[{i}] : {df[i].unique()}')
```

```
df['gender'] : ['Male' 'Female']
df['country'] : ['Sweden' 'Netherlands' 'Hungary' 'Belgium' 'Luxembourg' 'Italy'
'Croatia'
'Denmark' 'Malta' 'Germany' 'Poland' 'Ireland' 'Romania' 'Spain' 'Greece'
'Estonia' 'Cyprus' 'France' 'Slovenia' 'Latvia' 'Portugal' 'Austria'
'Czech Republic' 'Finland' 'Lithuania' 'Slovakia' 'Bulgaria']
df['cancer_stage'] : ['Stage I' 'Stage III' 'Stage IV' 'Stage II']
df['family_history'] : ['Yes' 'No']
df['smoking_status'] : ['Passive Smoker' 'Former Smoker' 'Never Smoked' 'Current
Smoker']
df['treatment_type'] : ['Chemotherapy' 'Surgery' 'Combined' 'Radiation']
```

# 1 Finish !

## 1.1 Deepak Kumar.

Thank you !